

Solution Brief: Knowledge Your AI Agents Can Actually Use

Turn scattered documents, data and relationships — across every platform you already run — into a governed knowledge layer, so every question gets a correct, current, cited answer.

THE CHALLENGE

Enterprise AI agents are failing in production because they face an Enterprise Context Gap. High-value decisions — approving a loan, diagnosing a machine, rerouting a supply chain — need both the live operational numbers in databases and the business rules locked in contracts, manuals and tickets. Those two worlds are disconnected, and the knowledge itself was built for humans, not machines: unreadable in document layouts, unconnected across systems, and unexplained in meaning, ownership and permissions. Agents therefore retrieve text that looks relevant rather than facts that are correct — unable to reason across documents, cite their sources, or tailor answers to who is asking. The bottleneck is no longer model capability. It is the missing governed knowledge layer between the data and the agent.

Gap	What it means	What it costs you
Unreachable	Facts sit in one platform, documents in another, across two or three clouds	No agent ever sees the whole picture — answers are right but incomplete
Unreadable	80–90% of knowledge is unstructured; standard extraction destroys tables and hierarchy	The agent quotes fragments with no idea what they belong to
Unconnected	The same customer appears under six names, with no record of how anything relates	Multi-hop and "what's affected?" questions are unanswerable at any model size
Ungoverned	Permissions aren't enforced at the point of retrieval	It can't go live on real data with real users
Unexplained	No agreed meaning, ownership or freshness for the underlying data	The agent picks the wrong source — and sounds completely certain

Why it's hard to solve	What good looks like
------------------------	----------------------

The bottleneck has moved from model capability to context. And that context is spread across platforms you've deliberately invested in — so the answer can't be "consolidate everything first." Getting the right knowledge, with the right permissions, into the model at the right moment is a data engineering and governance problem, solved across your estate rather than by rebuilding it.

Anyone — or any agent — asks a question in plain language and gets a correct, cited answer that respects their permissions. Including questions that span a contract, a customer record and a call transcript at the same time, held in three different systems.

THE SOLUTION

A governed **knowledge layer** between your data and your agents. It reads content that already exists — no one has to write anything new — and reaches across your existing platforms rather than replacing them. Three jobs, not one search box:

- **Understand** — read documents, calls and images properly, so structure, tables and meaning survive.
- **Connect** — resolve the six versions of a customer into one, and record how everything relates.
- **Serve** — answer with the right context, the right permissions and a citation you can check.

*Underlying components: borderless lakehouse, document understanding · embeddings · entity resolution · knowledge graph · hybrid GraphRAG retrieval · catalog-driven context

WHY GOOGLE CLOUD

Differentiator	Why it matters	Closes
Your data stays where it is	Open architectures and federated queries allow your AI to read facts held in existing third-party data warehouses - across any cloud - alongside your documents. Everything operates under one unified governance model with zero data migration or duplication required.	Unreachable
One engine for structured and unstructured	Our enterprise data platform natively stores both tables and documents, generates embeddings, runs vector search, and calls advanced LLMs all in one place. There is no need to stand up, secure, or reconcile a separate, bolt-on vector database.	Unreadable
Graph and vector in one consistent database	We execute complex relationship mapping (graph) and similarity search (vector) within a single, unified operational store. This eliminates the need to stitch two disjointed systems together - preventing the synchronization failures that derail most production GraphRAG deployments.	Unconnected
Governance you already run	A single enterprise business and metadata catalog and permission model spans all your structured and unstructured data providing agent context. Security policies and data lineage are strictly enforced at the moment of retrieval, rather than being patched on as a risky afterthought.	Ungoverned
Unified Semantic Metadata Layer	To automatically harvest metadata from disparate data sources and establish relationships between them. This unified context may include business definitions and operational lineage and provides freshness to the AI agents. By building a contextual model, with relationships, you dramatically improve the agent accuracy and reduce token cost.	Unexplained

OUTCOMES	PROOF
Time-to-answer — questions that take days answered in minutes	<i>[Insert nearest-fit named reference, metric and case study link]</i>
Groundedness — over 90% of answers supported by cited sources, measured continuously against a golden question set	Contract intelligence deployments on this pattern typically surface obligation and risk exposure in hours rather than weeks , with measured groundedness above 90%.
Reuse — each additional domain costs a fraction of the first, because the pipeline, governance and evaluation harness already exist	

GET STARTED

A two-hour discovery workshop. Bring your top 20 business questions; we'll map which ones your estate can answer today, where the gaps are, and which platforms hold the answers. Followed by an optional six-week proof of concept on one domain, with agreed success metrics.

Contact: [name · email] · **Learn more:** [link] · v1.2